


Project Name: AI Based Diabetes Prediction System

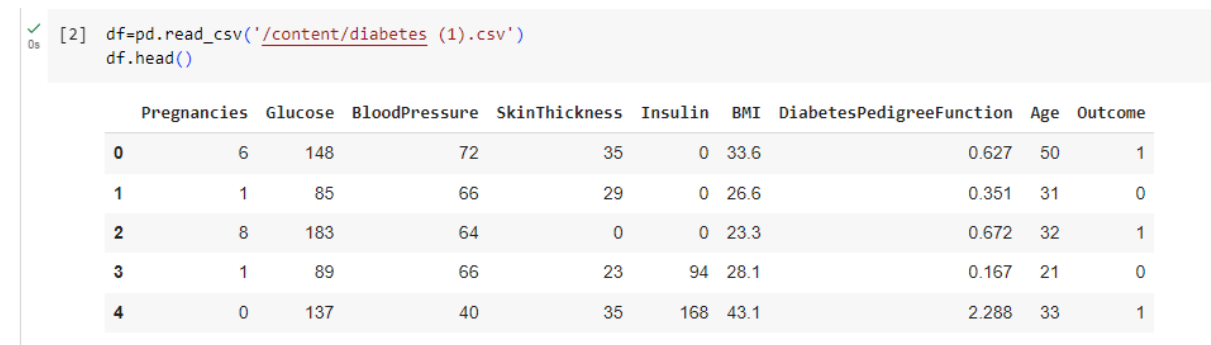
Project Code:203476

Coding:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```



```
df=pd.read_csv('/content/diabetes (1).csv')
df.head()
```

A screenshot of a Jupyter Notebook interface. The top part shows a code cell with the following code: `[2] df=pd.read_csv('/content/diabetes (1).csv')` and `df.head()`. Below the code cell, the output is displayed as a table with 10 columns: Pregnancies, Glucose, BloodPressure, SkinThickness, Insulin, BMI, DiabetesPedigreeFunction, Age, and Outcome. The table contains 5 rows of data, indexed from 0 to 4.

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
#outlier remove
```

```
Q1=df.quantile(0.25)
```

```
Q3=df.quantile(0.75)
```

```
IQR=Q3-Q1
```

```
print("----Q1--- \n",Q1)
```

```
print("\n---Q3--- \n",Q3)
```

```
print("\n---IQR---\n",IQR)
```

```
df_out = df[~((df < (Q1 - 1.5 * IQR)) |(df > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
df.shape,df_out.shape
```

```
X=df_out.drop(columns=['Outcome'])
```

```
y=df_out['Outcome']
```

```
#Splitting train test data 80 20 ratio
```

```
0s #outlier remove

Q1=df.quantile(0.25)
Q3=df.quantile(0.75)
IQR=Q3-Q1

print("----Q1--- \n",Q1)
print("\n---Q3--- \n",Q3)
print("\n---IQR---\n",IQR)
df_out = df[~((df < (Q1 - 1.5 * IQR)) |(df > (Q3 + 1.5 * IQR))).any(axis=1)]
df.shape,df_out.shape
X=df_out.drop(columns=['Outcome'])
y=df_out['Outcome']
#Splitting train test data 80 20 ratio
```

----Q1---	
Pregnancies	1.00000
Glucose	99.00000
BloodPressure	62.00000
SkinThickness	0.00000
Insulin	0.00000
BMI	27.30000
DiabetesPedigreeFunction	0.24375
Age	24.00000
Outcome	0.00000

Name: 0.25 dtype: float64

```

✓ 0s Outcome 0.00000
Name: 0.25, dtype: float64

---Q3---
Pregnancies 6.00000
Glucose 140.25000
BloodPressure 80.00000
SkinThickness 32.00000
Insulin 127.25000
BMI 36.60000
DiabetesPedigreeFunction 0.62625
Age 41.00000
Outcome 1.00000
Name: 0.75, dtype: float64

---IQR---
Pregnancies 5.00000
Glucose 41.25000
BloodPressure 18.00000
SkinThickness 32.00000
Insulin 127.25000
BMI 9.30000
DiabetesPedigreeFunction 0.3825
Age 17.00000
Outcome 1.00000
dtype: float64

```

```

from sklearn.model_selection import train_test_split
train_X,test_X,train_y,test_y=train_test_split(X,y,test_size=0.2)
train_X.shape,test_X.shape,train_y.shape,test_y.shape

```

```

✓ 0s [5] from sklearn.model_selection import train_test_split
      train_X,test_X,train_y,test_y=train_test_split(X,y,test_size=0.2)
      train_X.shape,test_X.shape,train_y.shape,test_y.shape

      ((511, 8), (128, 8), (511,), (128,))

```

```

from sklearn.metrics import confusion_matrix,accuracy_score,make_scorer
from sklearn.model_selection import cross_validate

```

```

def tn(y_true, y_pred): return confusion_matrix(y_true, y_pred)[0, 0]
def fp(y_true, y_pred): return confusion_matrix(y_true, y_pred)[0, 1]
def fn(y_true, y_pred): return confusion_matrix(y_true, y_pred)[1, 0]
def tp(y_true, y_pred): return confusion_matrix(y_true, y_pred)[1, 1]

```

```
#cross validation purpose
```

```
scoring = {'accuracy': make_scorer(accuracy_score), 'prec': 'precision'}
```

```
scoring = {'tp': make_scorer(tp), 'tn': make_scorer(tn),  
          'fp': make_scorer(fp), 'fn': make_scorer(fn)}
```

```
def display_result(result):
```

```
    print("TP: ", result['test_tp'])
```

```
    print("TN: ", result['test_tn'])
```

```
    print("FN: ", result['test_fn'])
```

```
    print("FP: ", result['test_fp'])
```

```
#Logistic Regression
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import roc_auc_score
```

```
acc=[]
```

```
roc=[]
```

```
clf=LogisticRegression()
```

```
clf.fit(train_X, train_y)
```

```
y_pred=clf.predict(test_X)
```

```
#find accuracy
```

```
ac=accuracy_score(test_y,y_pred)
```

```
acc.append(ac)
```

```
#find the ROC_AOC curve
```

```
rc=roc_auc_score(test_y,y_pred)
```

```
roc.append(rc)
```

```
print("\nAccuracy {0} ROC {1}".format(ac,rc))
```

```
#cross val score
```

```
result=cross_validate(clf,train_X,train_y,scoring=scoring,cv=10)
```

```
display_result(result)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Accuracy 0.8515625 ROC 0.7956821026282853
 TP: [3 8 10 12 7 11 8 9 9 9]
 TN: [30 30 28 30 25 33 31 31 34 34]
 FN: [14 9 7 5 10 6 8 7 7 7]
 FP: [5 4 6 4 9 1 4 4 1 1]
 /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

 Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

 Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

 Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

#Naive Bayes Theorem

#import library

from sklearn.naive_bayes import GaussianNB

clf=GaussianNB()

clf.fit(train_X,train_y)

y_pred=clf.predict(test_X)

#find accuracy

ac=accuracy_score(test_y,y_pred)

acc.append(ac)

#find the ROC_AOC curve

rc=roc_auc_score(test_y,y_pred)

roc.append(rc)

print("\nAccuracy {0} ROC {1}".format(ac,rc))

#cross val score

result=cross_validate(clf,train_X,train_y,scoring=scoring,cv=10)

display_result(result)

✓
0s



```
#Naive Bayes Theorem
#import library
from sklearn.naive_bayes import GaussianNB

clf=GaussianNB()
clf.fit(train_X,train_y)
y_pred=clf.predict(test_X)
#find accuracy
ac=accuracy_score(test_y,y_pred)
acc.append(ac)

#find the ROC_AOC curve
rc=roc_auc_score(test_y,y_pred)
roc.append(rc)
print("\nAccuracy {0} ROC {1}".format(ac,rc))

#cross val score
result=cross_validate(clf,train_X,train_y,scoring=scoring,cv=10)
display_result(result)
```



```
Accuracy 0.796875 ROC 0.7819072313454336
TP:  [10 11  8 10  7  8 10 11  7 11]
TN:  [32 26 26 32 28 28 31 31 31 27]
FN:  [7 5 8 6 9 8 6 5 9 5]
FP:  [3 9 9 3 7 7 4 4 4 8]
```