

# **Project Name: AI Based Diabetes Prediction System**

**Project Code:203476**

## **Problem Definition:**

The problem is to build an AI-powered diabetes prediction system that uses machine learning algorithms to analyze medical data and predict the likelihood of an individual developing diabetes. The system aims to provide early risk assessment and personalized preventive measures, allowing individuals to take proactive actions to manage their health.

## **Objective:**

The main objective of an AI-based Diabetes Prediction System is to help detect and predict diabetes in individuals. This system leverages artificial intelligence and machine learning techniques to analyze various data sources, such as medical records, patient history, and potentially even genetic information, to make predictions about a person's risk of developing diabetes. Here are some specific objectives of such a system:

- **Early Detection:** Identify individuals at risk of developing diabetes before they exhibit symptoms or receive an official diagnosis. Early detection can lead to early intervention and better management of the disease.
- **Risk Assessment:** Evaluate an individual's risk factors for diabetes, including lifestyle factors (e.g., diet, exercise), genetic predisposition, and medical history, to provide personalized risk assessments.
- **Preventive Measures:** Suggest preventive measures and lifestyle changes to individuals at risk, such as dietary recommendations, exercise routines, and regular monitoring, to help them reduce their risk of diabetes.
- **Treatment Planning:** For individuals already diagnosed with diabetes, assist healthcare professionals in creating personalized treatment plans based on the patient's specific needs and medical history.
- **Data Integration:** Integrate and analyze various types of data, including patient records, medical imaging, and wearable device data, to provide a comprehensive view of the patient's health and diabetes risk.

- **Patient Education:** Educate patients about diabetes, its management, and the importance of adhering to recommended lifestyle changes and treatment plans.
- **Healthcare Resource Allocation:** Help healthcare providers allocate their resources more efficiently by identifying high-risk individuals who may require more frequent monitoring and intervention
- **Research and Insights:** Generate insights from large datasets to improve our understanding of diabetes, its risk factors, and potential avenues for prevention and treatment.
- **Continuous Monitoring:** Enable continuous monitoring of patients' health data and provide real-time alerts to healthcare providers if a patient's risk level changes significantly.
- **Cost Reduction:** Potentially reduce healthcare costs associated with diabetes by preventing or mitigating complications through early intervention and management.

## Design:

To build AI based diabetics prediction model we need to install some packages and they are:

- NumPy
- Pandas
- Matplotlib
- Scikit Learn

Dataset is taken from <https://www.kaggle.com/datasets/mathchi/diabetes-data-set>.

Since the data in the dataset is the raw data, it needs to undergo the following stages:

### • **Data Collection:**

Data has to be collected from various sources to form a dataset. The dataset should contain some medical details or features such as glucose levels, ages, pressure, BMI etc.

### • **Data Preprocessing or cleaning and Visualization:**

Since the dataset is collected from various sources it may consist of duplicates, nullable and Irrelevant data. So, data needs to be cleaned, normalized, replacing the nullable values with standard values etc. Once the data is cleaned it can be used to prepare for training.

### • **Feature Selection:**

We need to select relevant features like glucose levels, ages, pressure, BMI etc. that can cause diabetics.

## **Model Selection:**

We are going to build the model using various algorithms like:

- Logistic Regression

## **Evaluation:**

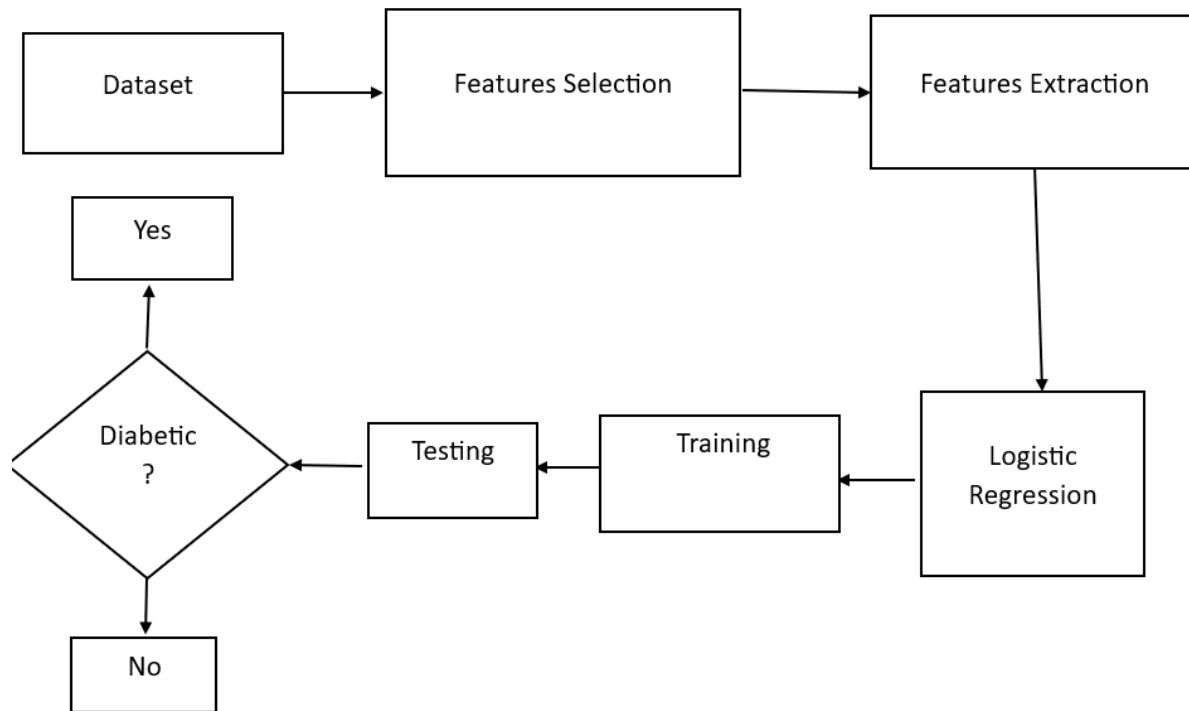
We will evaluate the model's performance using metrics like:

- Accuracy Score
- ROC AUC Curve
- Cross Validation
- Confusion Matrix

## **Iterative Improvement:**

We will fine-tune the model parameters and explore techniques like feature engineering to enhance prediction accuracy.

## **Innovative Design:**



## Coding:

#Importing packages

```
import pandas as pd
```

```
import numpy as np
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

✓  
1s

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

#Reading the dataset

```
df=pd.read_csv('/content/diabetes (1).csv')
```

```
df.head ()
```

```
df=pd.read_csv('/content/diabetes (1).csv')
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

df.describe ()

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

df.info()

```
[7] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Pregnancies           768 non-null   int64
 1   Glucose               768 non-null   int64
 2   BloodPressure         768 non-null   int64
 3   SkinThickness         768 non-null   int64
 4   Insulin               768 non-null   int64
 5   BMI                   768 non-null   float64
 6   DiabetesPedigreeFunction 768 non-null   float64
 7   Age                   768 non-null   int64
 8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

df.isnull ().values. any()

```
df.isnull().values.any()
```

False

Print (df.isnull ().sum ())

```
✓ 0s print(df.isnull().sum())
```

```
➔ Pregnancies      0  
   Glucose         0  
   BloodPressure   0  
   SkinThickness   0  
   Insulin         0  
   BMI            0  
   DiabetesPedigreeFunction  0  
   Age            0  
   Outcome         0  
   dtype: int64
```

df. Shape

```
✓ 0s [6] df.shape
```

```
(768, 9)
```


## Various Analysis

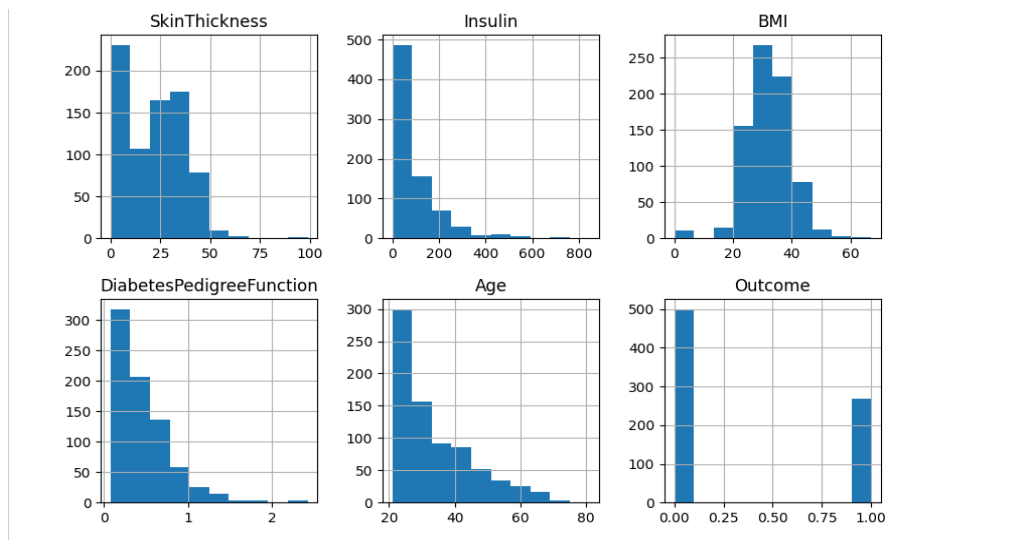
#Histogram

df.hist (bins=10, figsize= (10, 10))

plt.show ()

```
✓ 4s #histogram  
df.hist(bins=10,figsize=(10,10))  
plt.show()
```



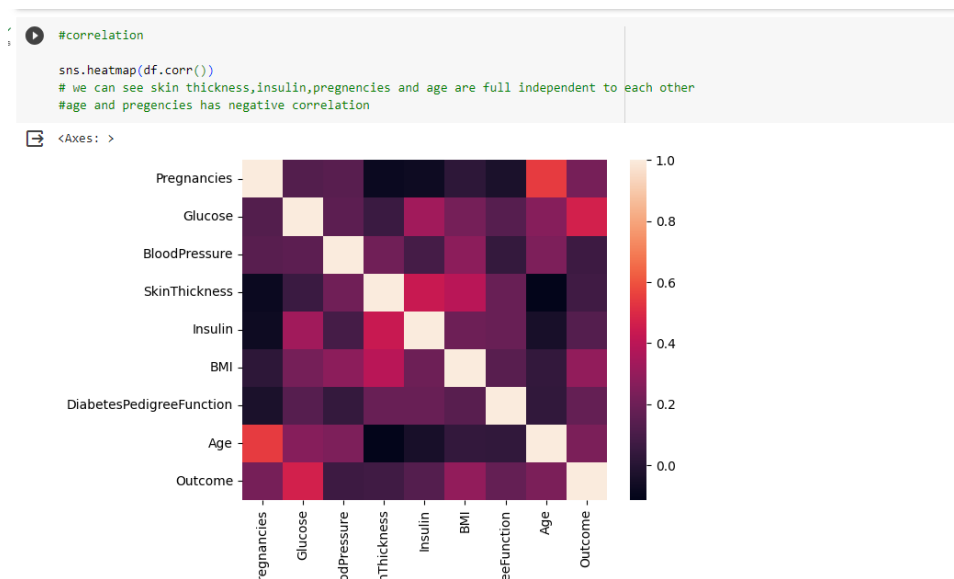


#correlation

sns.heatmap(df.corr())

# We can see skin thickness, insulin, pregnancies and age are full independent to each other

#age and pregnancies has negative correlation



#lets count total outcome in each target 0 1


#0 means no diabeted

#1 means patient with diabetes

sns.countplot(y=df['Outcome'],palette='Set1')

```
✓ 0s [6] #lets count total outcome in each target 0 1
      #0 means no diabeted
      #1 means patient with diabtes
      sns.countplot(y=df['Outcome'],palette='Set1')
```

<Axes: xlabel='count', ylabel='Outcome'>



Sns. Set (style="ticks")

sns.pairplot(df, hue="Outcome")


#box plot for outlier visualization

Sns. Set (style="whitegrid")

df.boxplot(figsize=(15,6))

```
✓ 35s [7] sns.set(style="ticks")
      sns.pairplot(df, hue="Outcome")
      #box plot for outlier visualization
      sns.set(style="whitegrid")
      df.boxplot(figsize=(15,6))
```

<Axes: xlabel='Age', ylabel='Density'>



#box plot


Sns. Set (style="whitegrid")



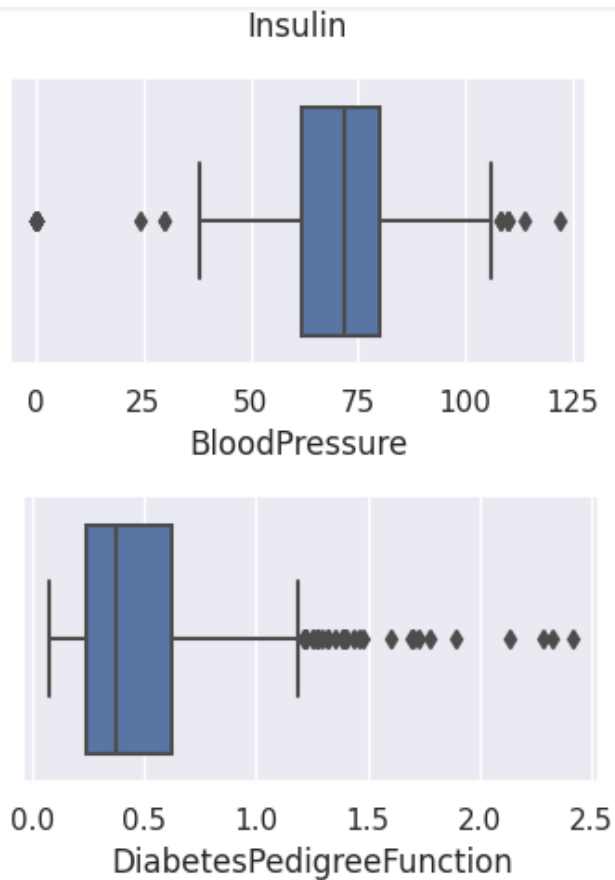
```
Sns. Set (arc= {'figure.figsize':(4,2)})
sns.boxplot(x=df ['Insulin'])
plt.show()
sns.boxplot(x=df['Blood Pressure'])
plt.show ()
sns.boxplot(x=df['DiabetesPedigreeFunction'])
plt.show ()
```

```
✓ 1s #box plot
sns.set(style="whitegrid")

sns.set(rc={'figure.figsize':(4,2)})
sns.boxplot(x=df['Insulin'])
plt.show()
sns.boxplot(x=df['BloodPressure'])
plt.show()
sns.boxplot(x=df['DiabetesPedigreeFunction'])
plt.show()
```



✓ [8]  
1s



```
[2] df=pd.read_csv('/content/diabetes (1).csv')  
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
Q1=df.quantile (0.25)
```

```
Q3=df.quantile (0.75)
```

```
IQR=Q3-Q1
```

```
print ("---#outlier remove -Q1--- \n", Q1)
```

```
print ("\n---Q3--- \n", Q3)
```

```
print("\n---IQR---\n",IQR)
```

```
df_out = df[~((df < (Q1 - 1.5 * IQR)) |(df > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
df.shape,df_out.shape
```

```
X=df_out.drop(columns=['Outcome'])
```

```
y=df_out['Outcome']
```

```
#Splitting train test data 80 20 ratio
```

```
#outlier remove

Q1=df.quantile(0.25)
Q3=df.quantile(0.75)
IQR=Q3-Q1

print("---Q1--- \n",Q1)
print("\n---Q3--- \n",Q3)
print("\n---IQR---\n",IQR)
df_out = df[~((df < (Q1 - 1.5 * IQR)) |(df > (Q3 + 1.5 * IQR))).any(axis=1)]
df.shape,df_out.shape
X=df_out.drop(columns=['Outcome'])
y=df_out['Outcome']
#Splitting train test data 80 20 ratio
```

```
---Q1---
Pregnancies      1.888889
Glucose           99.888889
BloodPressure     82.888889
SkinThickness     9.888889
Insulin           8.888889
BMI              27.388889
DiabetesPedigreeFunction  0.241775
Age              24.888889
Outcome           0.888889
```

```

Outcome      0.00000
Name: 0.25, dtype: float64

---Q3---
Pregnancies      6.00000
Glucose      140.25000
BloodPressure    80.00000
SkinThickness    32.00000
Insulin      127.25000
BMI      36.00000
DiabetesPedigreeFunction    0.62625
Age      41.00000
Outcome      1.00000
Name: 0.75, dtype: float64

---IQ3---
Pregnancies      5.00000
Glucose      41.25000
BloodPressure    18.00000
SkinThickness    32.00000
Insulin      127.25000
BMI      9.30000
DiabetesPedigreeFunction    0.3325
Age      17.00000
Outcome      1.00000
dtype: float64

```

```

from sklearn.model_selection import train_test_split
train_X,test_X,train_y,test_y=train_test_split(X,y,test_
size=0.)

train_X.shape,test_X.shape,train_y.shape,test_y.shape

```

```

[5] from sklearn.model_selection import train_test_split
train_X,test_X,train_y,test_y=train_test_split(X,y,test_size=0.2)
train_X.shape,test_X.shape,train_y.shape,test_y.shape

((511, 8), (128, 8), (511,), (128,))

```

```

from sklearn.metrics import
confusion_matrix,accuracy_score,make_scorer from
sklearn.model_selection import cross_validate

def tn(y_true, y_pred): return confusion_matrix(y_true,
y_pred)[0, 0] def fp(y_true, y_pred): return
confusion_matrix(y_true, y_pred)[0, 1] def fn(y_true,

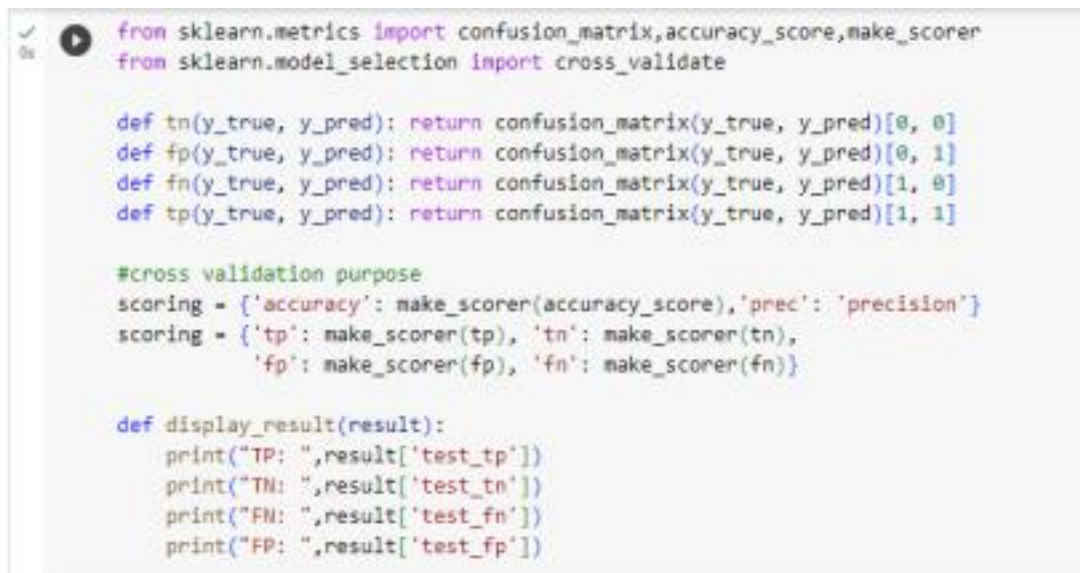
```

```

y_pred): return confusion_matrix(y_true, y_pred)[1, 0]
def
tp(y_true, y_pred): return confusion_matrix(y_true,
y_pred)[1, 1]

#cross validation purpose
scoring = {'accuracy':
make_scorer(accuracy_score), 'prec':

```



```

from sklearn.metrics import confusion_matrix, accuracy_score, make_scorer
from sklearn.model_selection import cross_validate

def tn(y_true, y_pred): return confusion_matrix(y_true, y_pred)[0, 0]
def fp(y_true, y_pred): return confusion_matrix(y_true, y_pred)[0, 1]
def fn(y_true, y_pred): return confusion_matrix(y_true, y_pred)[1, 0]
def tp(y_true, y_pred): return confusion_matrix(y_true, y_pred)[1, 1]

#cross validation purpose
scoring = {'accuracy': make_scorer(accuracy_score), 'prec': 'precision'}
scoring = {'tp': make_scorer(tp), 'tn': make_scorer(tn),
          'fp': make_scorer(fp), 'fn': make_scorer(fn)}

def display_result(result):
    print("TP: ", result['test_tp'])
    print("TN: ", result['test_tn'])
    print("FN: ", result['test_fn'])
    print("FP: ", result['test_fp'])

```

```

'Precision'} scoring = {'tp': make_scorer(tp), 'tn':
make_scorer(tn), 'fp': make_scorer(fp), 'fn': make_scorer
(fn)}

```

```

def display_result(result):
    Print("TP: ", result['testate'])
    Print ("TN: " result['testate'])
    Print ("FN: ", result['testify'])

```

```

print ("FP: ",result['precision'] scoring = {'tp':
make_scorer(tp), 'tn': make_scorer(tn), 'fp':
make_scorer(fp), 'fn': make_scorer(fn)}

```

```

def display_result(result):
    print("TP: ",result['test_tp'])
    print("TN: ",result['test_tn'])
    print("FN: ",result['test_fn'])
print("FP: ",result['test_fp'])'test_fp'])acc=[]
roc=[]

```

```

clf=LogisticRegression()
clf.fit(train_X,train_y)
y_pred=clf.predict(test_X)
#find accuracy
ac=accuracy_score(test_y,y_pred)

```

```

#Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score

acc=[]
roc=[]

clf=LogisticRegression()
clf.fit(train_X,train_y)
y_pred=clf.predict(test_X)
#find accuracy
ac=accuracy_score(test_y,y_pred)
acc.append(ac)

#Find the ROC_AUC
rc=roc_auc_score(test_y,y_pred)
roc.append(rc)
print("\nAccuracy {0} ROC {1}".format(ac,rc))

#cross val score
result=cross_validate(clf,train_X,train_y,scoring=scoring,cv=10)
display_result(result)

```

```
#find the ROC_AOC curve
rc=roc_auc_score(test_y,y_pred)
roc.append(rc)
print("\nAccuracy {0} ROC {1}".format(ac,rc))
```

```
#cross val score
result=cross_validate(clf,train_X,train_y,scoring=scoring,cv=10) display_result(result)
```

```
Accuracy: 0.813821 ROC: 0.798321899232813
TP: [ 2  0 18 12  7 11  8  9  0  0]
TN: [30 30 18 18 18 18 18 18 18 18]
FP: [34  0  7  5 30  0  8  7  7  7]
FN: [ 0  0  0  0  0  0  0  0  0  0]
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL no. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL no. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL no. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
#Naive Bayes Theorem
#import library
from sklearn.naive_bayes import GaussianNB
```

```
clf=GaussianNB()
clf.fit(train_X,train_y)
y_pred=clf.predict(test_X)
#find accuracy
ac=accuracy_score(test_y,y_pred)
acc.append(ac)
```

```

#find the ROC_AOC curve
rc=roc_auc_score(test_y,y_pred)
roc.append(rc)
print("\nAccuracy {0} ROC {1}".format(ac,rc))

#cross val score
result=cross_validate(clf,train_X,train_y,scoring=scori
ng,cv=10) display_result(result)

```



The screenshot shows a Jupyter Notebook cell with the following code:

```

#Naive Bayes Theorem
#Import library
from sklearn.naive_bayes import GaussianNB

clf=GaussianNB()
clf.fit(train_X,train_y)
y_pred=clf.predict(test_X)
#find accuracy
ac=accuracy_score(test_y,y_pred)
acc.append(ac)

#find the ROC_AOC curve
rc=roc_auc_score(test_y,y_pred)
roc.append(rc)
print("\nAccuracy {0} ROC {1}".format(ac,rc))

#cross val score
result=cross_validate(clf,train_X,train_y,scoring=scoring,cv=10)
display_result(result)

```

The output of the code is displayed below the code cell:

```

Accuracy 0.796875 ROC 0.7819072313454336
TP: [18 11  8 10  7  0 10 11  7 11]
TN: [32 26 26 32 28 28 31 31 31 27]
FN: [ 7  5  8  8  9  8  6  5  9  5]
FP: [ 3  9  9  3  7  7  4  4  4  8]

```





